

## 1. Scope

Build a minimal web or mobile app that lets the user:

- Sign in with email and password.
- Record voice journals.
- Store audio files and transcripts:
  - Locally in the app, grouped by calendar day.
  - In Google Drive (audio + transcript) if authorized.
- Use AI to merge and reason over multiple recordings from the same day.

No marketing site, no extra pages. Straight to "Journal" and sign in.

---

## 2. Authentication and Landing Flow

### Goals

- Simple sign-in.
- Single screen with "Journal" label, then sign-in form.

### Tasks

1. Implement email/password authentication.
  - Create sign up and sign in with:
    - Email
    - Password
  - Basic validation and error states.
  - Password reset flow if needed.
2. Landing page layout:
  - Header text: "Journal".
  - Subsection: sign-in / sign-up form directly on the same screen.
  - No extra navigation, no marketing copy, no hero section.
3. Post-login behavior:
  - On successful login, go directly to the "Journal" screen listing entries by day.
  - Persist session so user does not have to sign in every time.

---

## 3. Journal Screen and UX

### Goals

- Let user easily record a new entry.
- Show past entries grouped by day.
- Keep it utilitarian; function over aesthetics.

## Tasks

1. Journal main view:
    - o Top section: "Record new entry" with:
      - Record button (start / stop).
      - Recording timer indicator.
    - o Below: list of days.
      - Each day shows:
        - Date (YYYY-MM-DD or localized).
        - Summary snippet of that day (e.g. first line of transcript or AI generated short summary if available).
        - Expand/collapse to show individual recordings for that day if needed.
  2. Per-day drill-down:
    - o When a day is expanded:
      - Show list of individual recordings:
        - Timestamp of each recording.
        - Short preview of transcript.
        - Playback button for the audio.
      - Option to see the merged "day-level" transcript / summary produced by AI.
  3. Basic controls:
    - o Delete recording.
    - o Optional: rename a day or add a manual note.
- 

## 4. Audio Recording and Local Storage

### Goals

- Record audio from device mic.
- Store audio locally and tag it with correct date and time.

### Tasks

1. Implement microphone recording:
  - o Use appropriate API for platform (MediaRecorder or native mobile equivalent).
  - o Capture audio in a reasonably compressed format (e.g. WebM/Opus or AAC).
2. Associate metadata:
  - o For each recording, store:
    - id
    - user\_id
    - created\_at (timestamp)
    - local\_path or blob reference
    - duration
    - Link to transcript record (when available)
3. Local persistence:
  - o Implement a lightweight local database or storage:
    - Example: IndexedDB / SQLite / local file system depending on platform.
  - o Ensure recordings and transcripts are available offline.
4. Background edge cases:
  - o Handle cases where recording is interrupted or user navigates away.

---

## 5. Transcription

### Goals

- Convert recorded audio to text.
- Prefer on-device transcription if realistic.

### Tasks

1. Choose transcription approach:
    - Primary preference: on-device model (e.g. small Whisper variant or OS level speech-to-text) running locally.
    - Fallback: server-side transcription via API if on-device is not feasible. Make this configurable.
  2. Integrate transcription pipeline:
    - After recording stops:
      - Trigger transcription job.
      - Store transcript text in local storage associated with the audio.
    - Show loading state while transcription is in progress.
  3. Data model:
    - Transcript fields:
      - id
      - recording\_id
      - text
      - created\_at
      - Optional: language, confidence.
  4. Error handling:
    - If transcription fails:
      - Keep audio anyway.
      - Show error and offer "retry transcription".
- 

## 6. Google Drive Integration

### Goals

- Automatically upload each recording and its transcript to Google Drive if user connects Google.
- Keep a clean folder structure in Drive.

### Tasks

1. Google OAuth setup:
  - Add "Connect Google Drive" action in settings or journal screen.
  - Request appropriate Drive scopes with minimal permissions required.
  - Store and refresh tokens securely.
2. Drive file structure:
  - Create root folder, e.g. "JournalApp" (configurable).
  - For each day:
    - Optional: subfolder per date, e.g. JournalApp/2025-12-10/.
  - For each recording:

- Upload audio file with naming scheme like YYYY-MM-DD\_HH-mm-ss\_audio.ext.
- Upload transcript as .txt or .md with similar naming.

### 3. Sync logic:

- When a new recording + transcript is ready:
  - Queue an upload job.
  - Retry on transient failures.
  - Mark local entry as synced or sync\_failed.

### 4. Settings:

- Allow toggling Drive sync on/off.
  - Show sync status indicators per entry.
- 

## 7. AI Grouping and Day-Level Aggregation

### Goals

- Group multiple recordings from the same calendar day into a single logical "day entry".
- Provide AI features that operate over all recordings in that day and across days.

### Tasks

1. Day grouping logic:
  - Define grouping key: local calendar date based on created\_at and user timezone.
  - All recordings whose timestamps fall on the same date are grouped.
  - If user records twice in a day or multiple days in one sitting:
    - Use timestamp of when the recording is created to assign to the correct date.
  - Handle edge cases around midnight and timezone.
2. Day-level aggregation:
  - For each day:
    - Merge transcripts of all recordings for that day in chronological order.
  - Save this merged transcript as a separate entity:
    - day\_id, date, full\_text.
3. AI processing per day:
  - Run an AI pass on the merged day transcript to:
    - Generate a short summary.
    - Optional: extract key themes, tasks, moods.
4. AI across multiple days:
  - Provide an interface to query over multiple days:
    - Example use cases:
      - "Show me themes from last week."
      - "What was I worried about most this month."
  - Backend or on-device vector search:
    - Index day-level transcripts and summaries.
    - Answer queries by retrieving relevant days and constructing responses.
5. Local-first preference:
  - If using an AI model:
    - Prefer running small models on device if the platform allows.
    - If using a cloud model, keep PII and raw audio exposure minimal and document that.

## 8. Local-first and Privacy Constraints

### Goals

- Maximize on-device processing and storage.
- Minimize server dependencies.

### Tasks

1. Architecture choices:
  - Keep user data (audio + transcript) stored locally by default.
  - Only send data off-device for:
    - Explicit Drive sync.
    - AI or transcription cloud services if on-device is not available.
2. Configurable cloud usage:
  - Settings to toggle:
    - "Use cloud transcription" on/off.
    - "Use cloud AI" on/off.
  - Clear warnings when data is sent off-device.
3. Documentation for user:
  - Simple explanation somewhere in-app:
    - Where data is stored.
    - When it goes to Google Drive.
    - When, if ever, it goes to external AI/transcription APIs.

---

## 9. Remove Unnecessary Homepage

### Goals

- No marketing or landing homepage.

### Tasks

1. Remove existing homepage route and assets.
  2. Set application root route to the "Journal" sign-in screen.
  3. Ensure navigation after login goes directly to journal list.
-